
COMP30023 – Computer
Systems

Socket Programming and TCP flow control

Dr Lachlan Andrew



Recap

- TCP
 - Connection
 - Sending Data
 - Disconnection

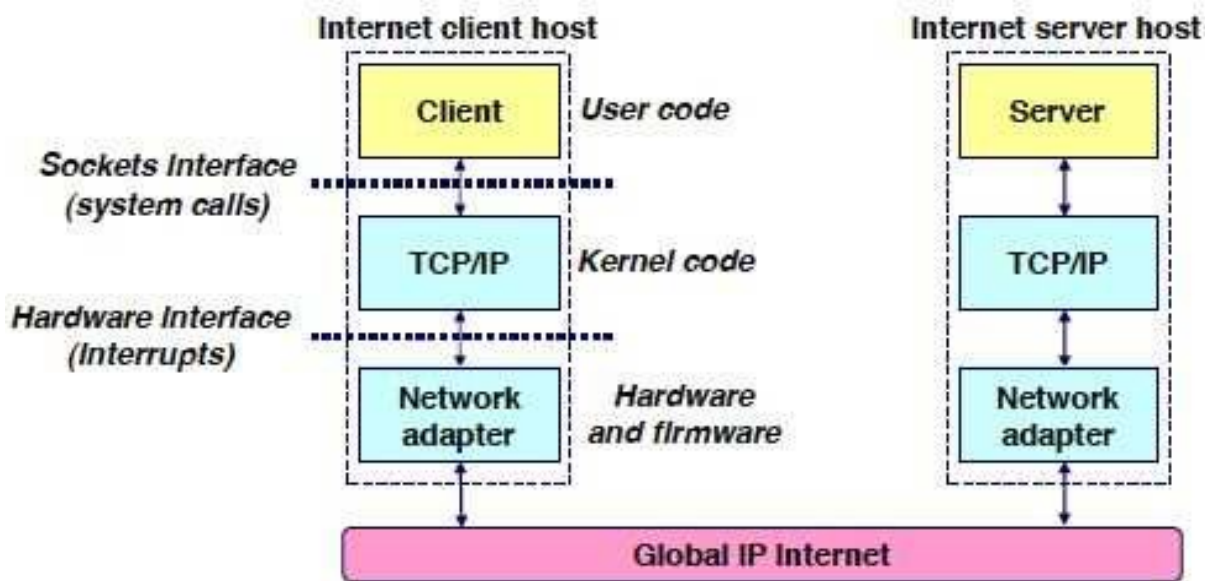


Summary

- High level overview of socket programming
 - Relationship to TCP
 - Sockets in C

Sockets

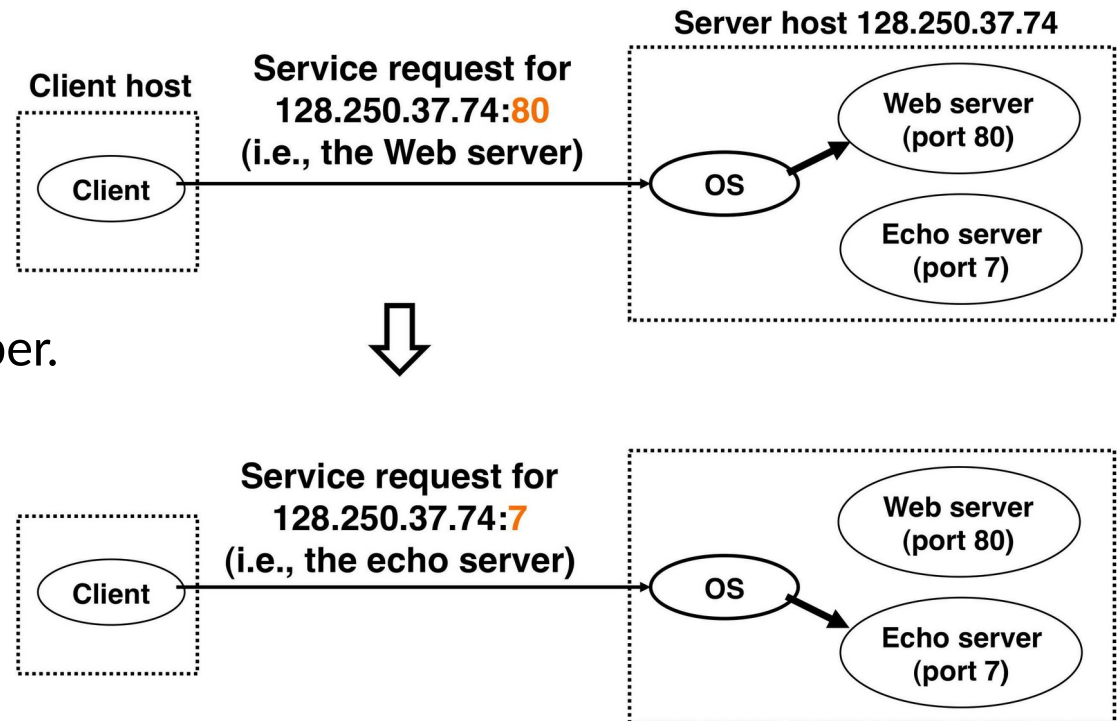
- A message going from one host to another must cross the underlying network.
- A process sends and receives through a socket
 - the “doorway” leading in/out of the application



Using sockets

- The “address” of a socket is the 5-tuple:

- Protocol
- source-IP
- source-port number
- destination-IP
- destination-port number.

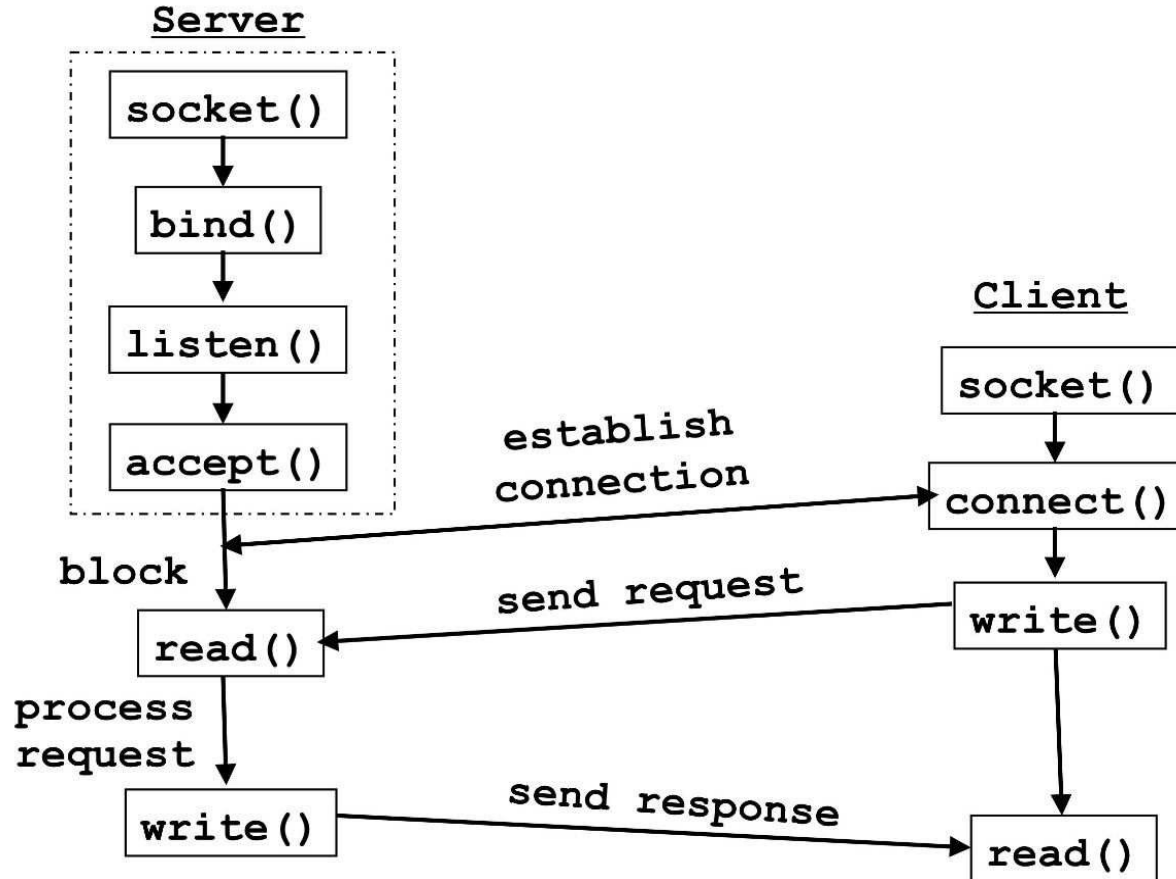




Berkeley Sockets

- Socket interface
 - originally provided in Berkeley UNIX
 - later adopted by all popular operating systems
 - simplifies porting applications to different OSes
- In UNIX, everything is like a file
 - all input is like reading a file
 - all output is like writing a file
 - file is “addressed” by an integer file descriptor
- API implemented as system calls:
 - examples include `connect()`, `read()`, `write()`, `close()`

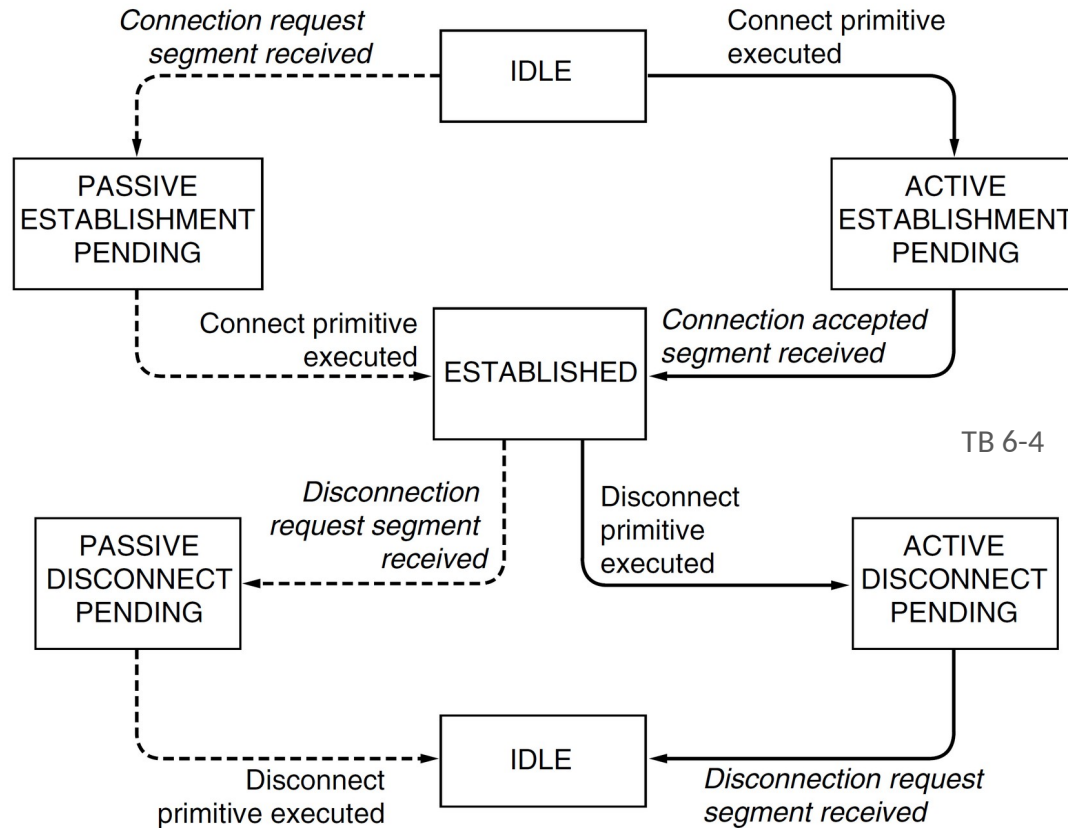
Using sockets



Socket Primitives

| State | Description |
|---------|---|
| SOCKET | Creates a new communication endpoint |
| BIND | Associate a local address with a socket |
| LISTEN | Announce willingness to accept connections; give queue size |
| ACCEPT | Passively establish an incoming connection (block until then) |
| CONNECT | Actively attempt to establish a connection |
| SEND | Send some data over a connection (write()) |
| RECEIVE | Receive <i>some</i> data from the connection (read()) |
| CLOSE | Release the connection |

Simplified (6-)State diagram for Connection Management



Complete list of Socket States

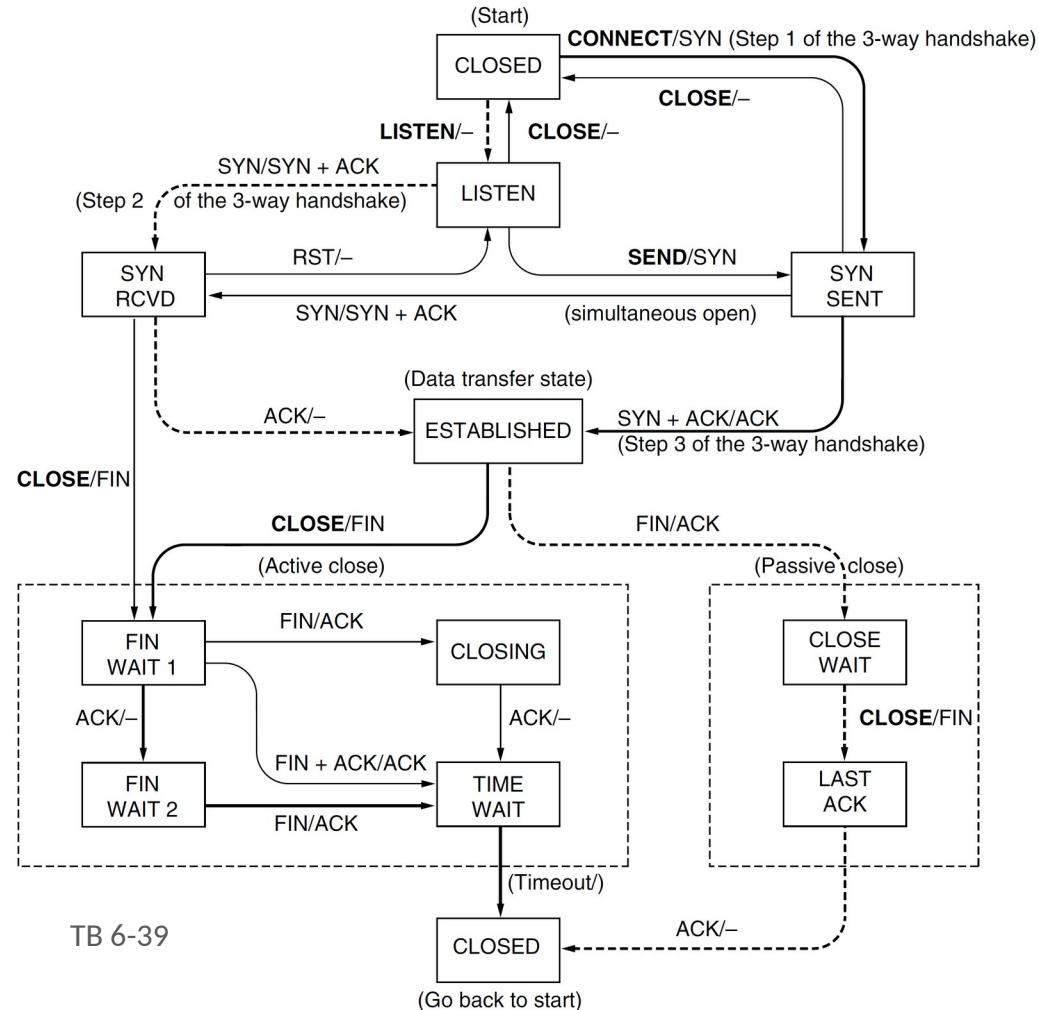
| State | Simplified name | Description |
|-------------|-----------------|--|
| CLOSED | Idle | No connection is active or pending |
| LISTEN | Pass. est. | The server is waiting for an incoming call |
| SYN RCVD | Pass. est. | A connection request has arrived; wait for ACK |
| SYN SENT | Act. est. | The application has started to open a connection |
| ESTABLISHED | Established | The normal data transfer state |
| FIN WAIT 1 | Act. disc. | The application has said it is finished |
| FIN WAIT 2 | Act. disc. | The other side has agreed to release |
| TIME WAIT | Act. disc. | Wait for all packets to die off |
| CLOSING | Act. disc. | Both sides have tried to close simultaneously |
| CLOSE WAIT | Pass. disc. | The other side has initiated a release |
| LAST ACK | Pass. disc. | Wait for all packets to die off |

Socket Finite State Machine

Bold before slash:
System call
e.g., connect

Non-bold before
slash:
Packet received
e.g., SYN

After slash:
Packet sent
e.g., SYN





Sockets in C

- Headers

```
#define _POSIX_C_SOURCE 200112L // Required for VSCode
#include <netdb.h>
#include <string.h>
#include <stdio.h>
#include <unistd.h>
```

- Variables

```
int listenfd = 0, connfd = 0, re = 1, s, n;
char sendBuff[1024];
struct addrinfo hints, *res, *rp;
```

- Create a socket

```
memset(&hints, 0, sizeof hints);
hints.ai_family = AF_INET;
hints.ai_socktype = SOCK_STREAM;
hints.ai_flags = AI_PASSIVE;
s = getaddrinfo(NULL, "5000", &hints, &res);
listenfd = socket(res->ai_family, res->ai_socktype, res->ai_protocol);
setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR, &re, sizeof(re));
```

AF_INET6

Create IPv6 socket

- getaddrinfo() can return multiple addresses as a linked list
- If you want the IPv6 address in particular, you may have to loop over the responses:

```
for (p = res; p != NULL; p = p->ai_next) {  
    if (p->ai_family == AF_INET6  
        && (sock = socket(p->ai_family,  
                          p->ai_socktype,  
                          p->ai_protocol)) < 0 )  
    {  
        // socket creation was attempted but failed  
    }  
}  
// Use sock here if it is >= 0.
```

Sockets in C - server

- Bind and listen

```
bind(listenfd, res->ai_addr, res->ai_addrlen);  
// maximum number of client connections to queue  
listen(listenfd, 10);
```

p of loop

Put TCP state machine in LISTEN state

- Accept, write/send, close

```
struct sockaddr_storage client_addr;  
socklen_t client_addr_size = sizeof client_addr;  
connfd = accept(listenfd, (struct sockaddr*)&client_addr,  
&client_addr_size);  
snprintf(sendBuff, sizeof(sendBuff), "Hello World!\n");  
n = write(connfd, sendBuff, strlen(sendBuff));  
close(connfd);
```

Block

- Wait until one of several files is ready to read / write
 - select(), pselect(), poll()

Sockets in C - client

- Connect

```
// Same as server, without hints.ai_flags
s = getaddrinfo("127.0.0.1", "5000", &hints, &res);
for (rp = res; rp != NULL; rp = rp->ai_next) {
    connfd = socket(rp->ai_family, rp->ai_socktype, rp->ai_protocol);
    if (connfd == -1) continue;
    if (connect(connfd, rp->ai_addr, rp->ai_addrlen) != -1) break;
    close(connfd);
}
```

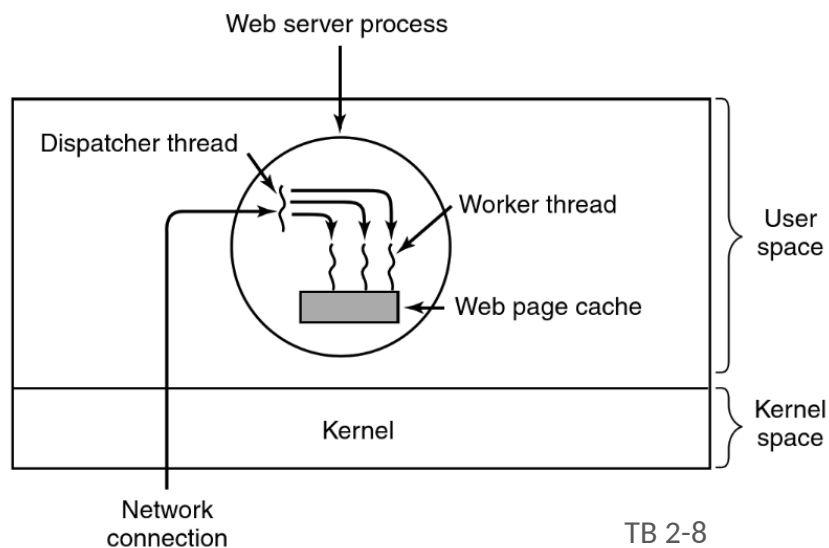
- Receive

```
while ((n = read(connfd, recvBuff, sizeof(recvBuff)-1)) > 0) {
    // process received buffer
}
```

- If the socket is blocking (see `fcntl` and `O_NONBLOCK`), it waits until there is data
 - This loop reads the whole connection
- If non-blocking, this just reads data that has arrived
 - More may come after a delay

Multi-threaded Web Server

- Clearly a web server needs to be able to handle concurrent connections from multiple clients
- This can be achieved through the usage of a multi-threaded web server



Multi-threaded web server

```
while (TRUE) {  
    get_next_request(&buf);  
    handoff_work(&buf);  
}
```

(a)

```
while (TRUE) {  
    wait_for_work(&buf)  
    look_for_page_in_cache(&buf, &page);  
    if (page_not_in_cache(&page))  
        read_page_from_disk(&buf, &page);  
    return_page(&page);  
}
```

TB 2-9

(b)

- High level outline of code for previous slide:
 - Dispatcher thread
 - Worker thread

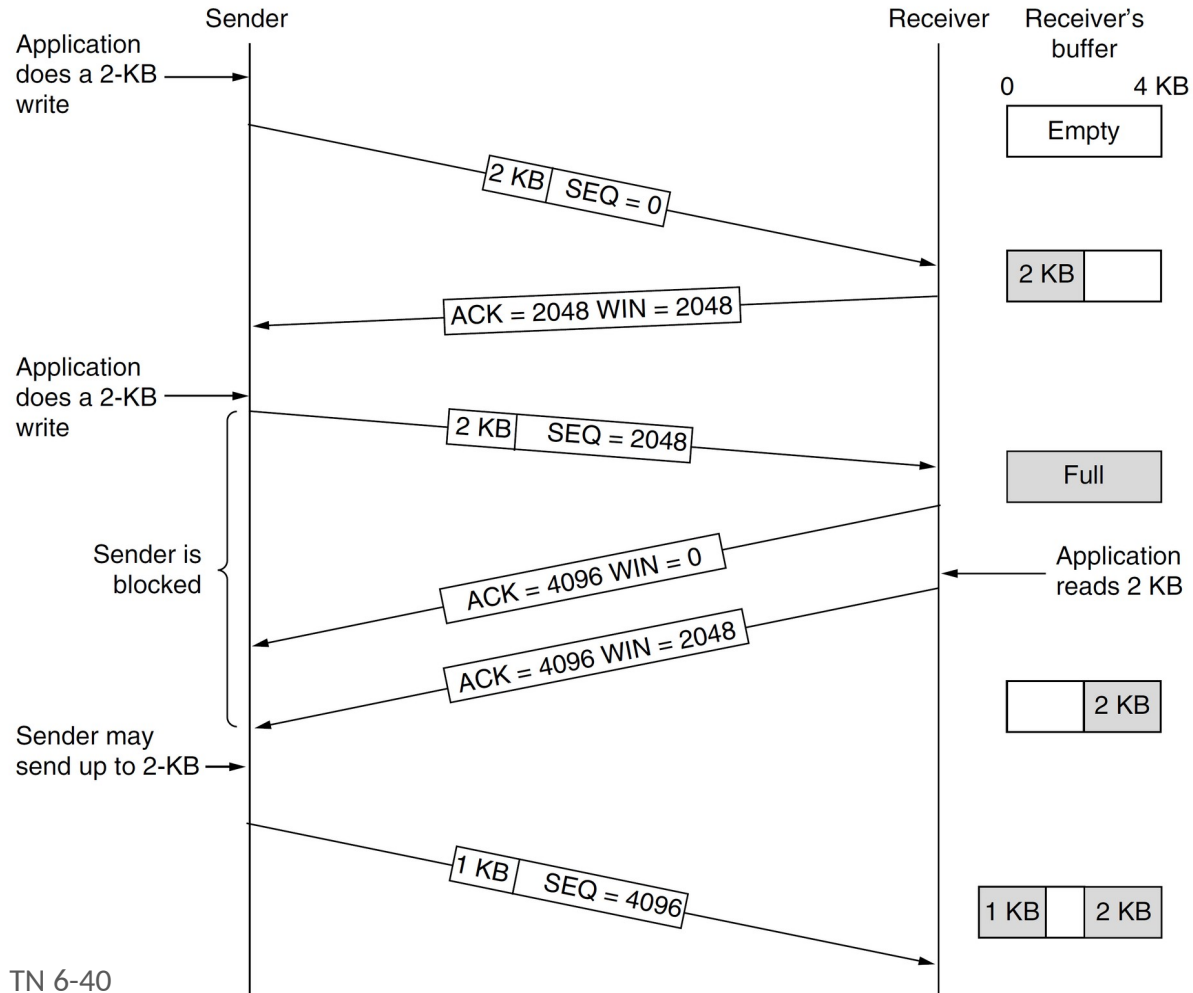




TCP Sliding Window

- Sliding window is controlled by receiver
- Determines amount of data the receiver is able to accept
 - Sender and receiver maintain buffers to send and receive data independently of the application
 - No guarantee that data is immediately sent or read from the respective buffers

TCP Sliding Window



TN 6-40

TCP Sliding Window

- When the window is 0 the sender should not send any data
 - Can send URGENT data
 - Can send “zero window probe”: 0 byte segment that causes the receiver to re-announce the next expected byte and window size (window probe) this is designed to prevent deadlock
- Senders may delay sending data, e.g., instead of sending the 2kiB immediately, could wait for a further 2kiB to fill the 4kiB receive window




TCP Sliding Window

- Send window
 - What data the sender is able to send – unacknowledged segments and unsent data that will fit into the receive window
- Receive window
 - Amount of data the receiver is willing to receive – window size in ACK
- Other windows are maintained for congestion control

TCP Sliding Window

| | | | | | | | | | | | | | | |
|---|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|-----|
| 1 | 11 | 21 | 31 | 41 | 51 | 61 | 71 | 81 | 91 | 101 | 111 | 121 | 131 | 141 |
|---|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|-----|

Send Window



Sender


Segment Size 10 bytes
SYN-SYN/ACK-ACK Completed
SYN:1, ACK:1, Window:50

Receive Window




Receiver


TCP Sliding Window




TCP Sliding Window




TCP Sliding Window




TCP Sliding Window




TCP Sliding Window




TCP Sliding Window




TCP Sliding Window




TCP Sliding Window



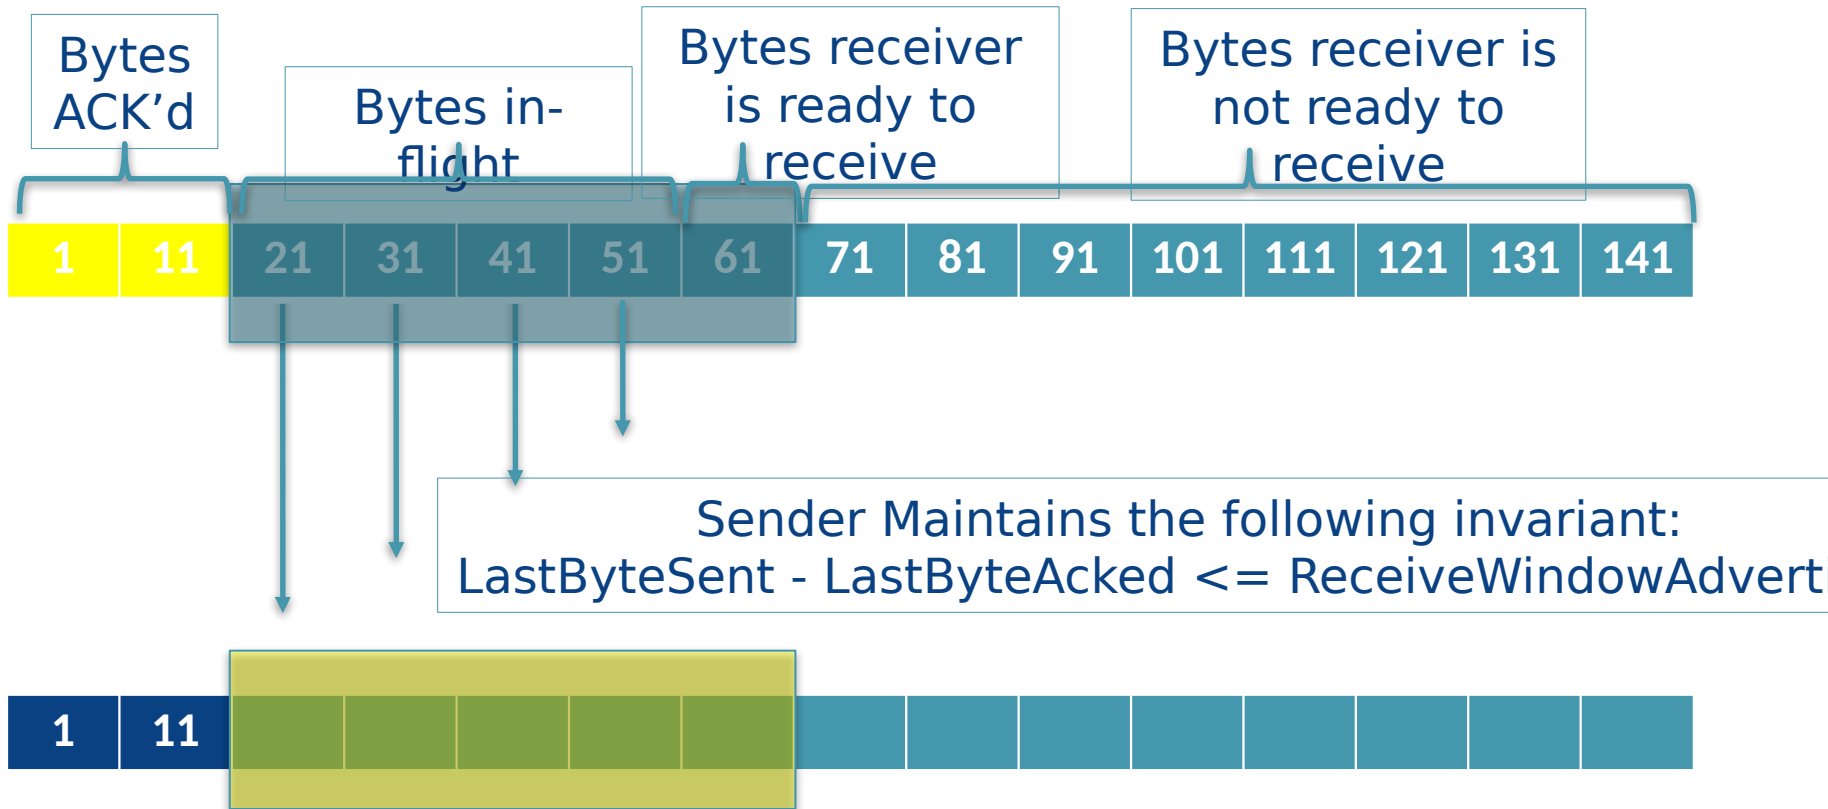
TCP Sliding Window – Window Update



TCP Sliding Window



TCP Sliding Window



Acknowledgement

- The slides are based on slides prepared by Chris Culnane based on material developed previously by: Michael Kirley, Zoltan Somogyi, Rao Kotagiri, James Bailey and Chris Leckie.
- Some of the images included in the notes were supplied as part of the teaching resources accompanying the text books listed in lecture 1.
 - (And also) Computer Networks, 6th Edition, Tanenbaum A., Wetherall. D.
<https://ebookcentral.proquest.com/lib/unimelb/detail.action?docID=6481879>
- Textbook Reference: 3.5.4, 3.5.5, 3.4, bits of 3.2. The text doesn't cover sockets in C, but many books do; Google's first pick is: TCP/IP Sockets in C, by Donahoo. M, Calvert. K